Intelligent Water Distribution Monitoring and Leakage Detection Platform

A CAPSTONE PROJECT REPORT

Submitted in the partial fulfilment for the Course of

CSA1016 - Software Engineering

to the award of the degree of

BACHELOR OF ENGINEERING
IN
Information Technology

Submitted by
192321050 - LENIN K

Under the Supervision of
Dr. K Anita Davamani

SIMATS ENGINEERING
Saveetha Institute of Medical and Technical Sciences
Chennai-602105

August 2026

SIMATS ENGINEERING

Saveetha Institute of Medical and Technical

Sciences Chennai-602105

August 2025

SIMATS ENGINEERING
Saveetha Institute of Medical and Technical Sciences
Chennai-602105

DECLARATION

I, Lenin K, Register Number 192321050, of the Department of Information Technology, SIMATS Engineering, Saveetha Institute of Medical and Technical Sciences, Saveetha University, Chennai, hereby declare that the Capstone Project Work entitled “Intelligent Water Distribution Monitoring and Leakage Detection Platform” is the result of my own bonafide efforts. To the best of my knowledge, the work presented herein is original, accurate, and has been carried out in accordance with principles of engineering ethics.

Place: Chennai
Date:

Signature of the Student

ABSTRACT

Water is among the most critical resources managed by urban infrastructure, yet distribution networks worldwide continue to lose a substantial share of treated water before it reaches consumers. Aging pipelines, undetected leaks, unauthorized consumption, and delayed maintenance responses contribute to this loss, commonly referred to as non-revenue water. Conventional monitoring practices rely heavily on periodic manual inspection, customer complaints, and billing reconciliation, all of which identify problems only after significant water and revenue have already been lost. This capstone project addresses the resulting gap by proposing an Intelligent Water Distribution Monitoring and Leakage Detection Platform intended to bring continuous, data-driven visibility to urban water networks within an academic prototype context.

The proposed system ingests real-time readings from IoT-enabled pressure sensors, flow meters, water quality sensors, and smart valves distributed across a water distribution network. Using statistical anomaly detection and predictive analytics, the platform identifies pressure and flow irregularities consistent with leakage, flags abnormal consumption patterns that may indicate unauthorized use, and forecasts pipeline segments at elevated risk of failure based on historical and sensor-derived indicators. Detected anomalies trigger prompt notifications to maintenance teams, while an operational dashboard consolidates network status, water quality trends, and consumption analytics for utility managers. Role-based access separates the responsibilities of administrators, control-room operators, and field maintenance staff, and the architecture is designed for secure, fault-tolerant, city-wide deployment built on established IoT messaging and time-series storage technologies rather than a custom-built telemetry stack.

The report documents the requirements analysis, system architecture, design decisions, and evaluation methodology underlying the platform's development. Expected outcomes include a functional prototype demonstrating reliable sensor data ingestion, timely leak detection, and actionable maintenance alerts, alongside an assessment of the system's strengths and limitations as an academic capstone project.

TABLE OF CONTENTS

CHAPTER 1: INTRODUCTION

1.1 Introduction

Water distribution has traditionally been managed through infrastructure designed decades ago, monitored primarily through periodic manual readings and reactive responses to customer complaints. As cities grow and water scarcity intensifies in many regions, utility providers face increasing pressure to reduce losses, extend the operational life of ageing infrastructure, and respond to faults before they escalate into service disruptions. Despite this pressure, the tools available to most utilities for day-to-day distribution monitoring have not kept pace with the sensing, connectivity, and analytics capabilities that are now routine in other infrastructure domains such as electricity grids and transportation networks.

This chapter introduces the capstone project, an Internet of Things (IoT) based platform designed to continuously monitor pressure, flow, consumption, and water quality across a distribution network, detect leakage and abnormal conditions early, and support proactive maintenance decisions. The chapter establishes the background and industry context motivating the project, identifies the problem being addressed, and defines the aim, objectives, and scope guiding the remainder of the report.

1.2 Background

Water loss in distribution networks, generally termed non-revenue water (NRW), includes both physical losses from leaks and bursts and apparent losses from metering inaccuracies and unauthorized consumption. Utilities in many regions report that a substantial proportion of treated water is lost before billing, representing both an economic burden and an avoidable strain on already scarce water resources. Historically, leak detection depended on visible surface indicators, acoustic listening devices operated manually along pipe routes, or customer-reported disruptions, all of which tend to identify leaks only after a considerable volume of water has already been lost.

The broader trajectory of infrastructure monitoring has shifted toward continuous, sensor-based approaches. Supervisory Control and Data Acquisition (SCADA) systems introduced centralized monitoring of key network points such as pumping stations and reservoirs, but conventional SCADA deployments typically instrument only a limited number of critical assets rather than the distribution network as a whole, owing to the cost of wired instrumentation. The maturation of low-cost wireless sensors, low-power wide-area networking, and cloud-based analytics has made it increasingly feasible to instrument distribution networks far more densely than SCADA alone allows, enabling detection approaches based on continuous pressure and flow analysis rather than point-in-time readings.

1.3 Industry Context

General-purpose IoT platforms and messaging frameworks, including MQTT-based brokers, time-series databases, and open-source dashboards such as ThingsBoard and Grafana, have made it considerably easier for organisations to build sensor-driven monitoring systems without developing telemetry infrastructure from first principles. In parallel, research into IoT-based water leak detection has grown substantially, exploring approaches ranging from low-cost water-level and flow sensing to deep-learning models trained on pressure and acoustic signals for both leak detection and localisation.

At the same time, adoption within municipal water utilities has been comparatively slow. Many utilities continue to rely on SCADA systems designed around a small number of monitored points, supplemented by manual patrols and customer complaints for leak identification across the wider distribution network. This creates a gap between what IoT-based sensing and analytics can technically achieve and what is actually deployed across most distribution networks, particularly at the scale of individual District Metered Areas (DMAs) rather than only trunk mains and treatment facilities.

Table 1.1 summarises the general progression of water network monitoring approaches, from manual inspection to fully instrumented, analytics-driven networks, to situate the proposed project within this broader evolution.

Table 1.1: Evolution of Water Distribution Monitoring Approaches

Stage	Approach	Representative Method	Monitoring Model
1	Manual inspection	Physical patrols, visible leak reporting	Reactive, point-in-time
2	Complaint-driven response	Customer complaints, billing reconciliation	Reactive, delayed
3	Centralized SCADA	Wired sensors at key assets (pumps, reservoirs)	Continuous but sparse
4	Smart metering (AMI)	Automated meter reading at consumer endpoints	Continuous, consumption-focused
5	Distributed IoT monitoring (proposed)	Wireless sensors network-wide with predictive analytics	Continuous, network-wide, predictive

As Table 1.1 illustrates, most utilities operate between Stage 2 and Stage 4, while distributed, analytics-driven monitoring across the full distribution network remains comparatively rare, representing the gap this project addresses.

1.4 Need for Intelligent Water Distribution Monitoring

Water utility operations involve several distinct roles, each requiring different information and levels of access to network data: field sensors continuously observe pipeline conditions, an analytics platform interprets that data, an alert engine escalates anomalies, and a maintenance team resolves faults in the field. Figure 1.1 illustrates a simplified representation of this workflow.

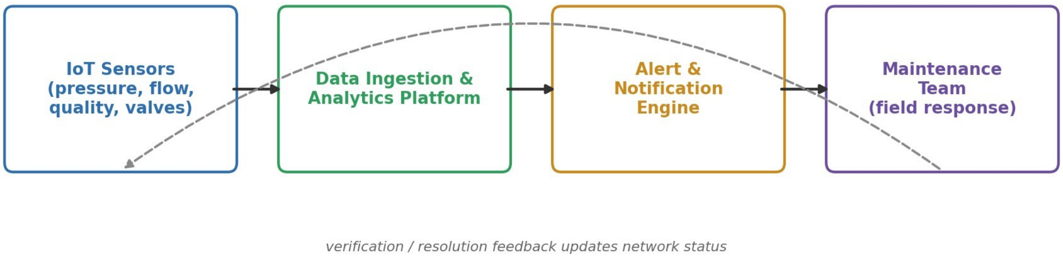

Figure 1.1: Simplified data and response workflow for the proposed monitoring platform, from field sensing through to maintenance resolution.

This workflow highlights a structural requirement that generic monitoring tools do not directly support: the need for dense, continuous sensing across the distribution network rather than only at a handful of critical assets, and the need for detected anomalies to be routed automatically to the appropriate maintenance response rather than depending on manual reporting.

Table 1.2 outlines the limitations of existing approaches to water distribution monitoring when applied to this workflow.

Table 1.2: Limitations of Existing Approaches to Water Distribution Monitoring

Existing Approach	Strength	Limitation for Water Distribution Monitoring
Traditional SCADA (key assets only)	Reliable, real-time at instrumented points	Sparse coverage; blind to leaks between monitored points
Smart metering / AMI	Accurate, detailed consumption data	Detects loss only retrospectively; no mid-network pressure/flow visibility
Manual leak survey	Can precisely locate confirmed leaks	Labour-intensive, infrequent, cannot scale city-wide
Customer complaint-based detection	Simple, no infrastructure investment	Delayed; only after substantial loss or visible damage
Generic IoT dashboards	Flexible; supports many sensor types	No built-in leak logic, pipeline topology, or water-domain analytics

As shown in Table 1.2, no single existing category of tool combines dense network-wide sensing with domain-specific leak detection, predictive maintenance, and role-based operational workflows. This gap forms the basis for the problem addressed by this project.

1.5 Problem Overview

Urban water distribution systems experience significant water losses due to pipeline leakages, unauthorized consumption, aging infrastructure, and delayed maintenance. Water utilities often identify leaks only after customer complaints or substantial water loss has occurred. With increasing water scarcity worldwide, utility providers require intelligent monitoring systems capable of detecting leakages early and optimizing water distribution efficiency.

This project addresses that requirement by designing and developing a Smart Water Distribution Monitoring Platform that continuously monitors water pipelines using IoT pressure sensors, flow meters, water quality sensors, and smart valves. The platform ingests real-time pressure, flow rate, consumption, and pipeline health data from the water distribution network, and, using predictive analytics, detects leakages, identifies abnormal consumption patterns, predicts infrastructure failures, notifies maintenance teams immediately, and generates comprehensive operational reports. The platform is designed to support city-wide deployment with secure communication, fault tolerance, and high availability.

1.6 Aim of the Project

The aim of this project is to design and implement an IoT-enabled water distribution monitoring platform that continuously ingests pressure, flow, consumption, and water quality data, applies predictive analytics to detect leakages and abnormal consumption patterns, forecasts infrastructure failures, and supports proactive maintenance and city-wide operational decision-making, using established IoT messaging and data-processing frameworks rather than a custom-built telemetry stack.

1.7 Objectives of the Project

The specific objectives of the project are as follows:

1. To implement continuous ingestion of pressure, flow rate, consumption, and pipeline health data from a distributed network of IoT sensors.
2. To design a leak detection mechanism based on statistical and machine-learning analysis of pressure and flow anomalies.
3. To implement water quality monitoring covering parameters relevant to distribution safety.
4. To develop predictive models capable of estimating the likelihood of infrastructure failure for individual pipeline segments.
5. To implement an alerting and notification mechanism that informs maintenance teams promptly when anomalies are detected.
6. To analyse consumption patterns to identify abnormal usage indicative of unauthorized tapping or metering faults.
7. To design an operational dashboard summarising network status, leak events, and maintenance activity for utility managers.
8. To implement role-based access control distinguishing between administrators, control-room operators, and field maintenance staff.
9. To design the platform architecture for secure communication, fault tolerance, and high availability suited to city-wide deployment.
10. To generate operational reports summarising water loss, maintenance response times, and conservation outcomes.

Table 1.3 maps each objective to the corresponding problem identified in Section 1.5, clarifying the rationale behind each planned feature.

Table 1.3: Mapping of Project Objectives to Identified Problems

Objective	Related Problem
Continuous sensor data ingestion	Absence of real-time visibility into network conditions
Leak detection via pressure/flow anomaly analysis	Delayed leak identification via complaints or manual survey
Water quality monitoring	No integrated visibility into water safety alongside loss monitoring
Predictive infrastructure failure modelling	Reactive maintenance on ageing infrastructure
Alerting and notification mechanism	Slow communication between detection and maintenance response
Consumption pattern analysis	Undetected unauthorized consumption and metering faults
Operational dashboard	Fragmented visibility across the distribution network
Role-based access control	Undifferentiated access for operators, administrators, and field staff
Secure, fault-tolerant, city-wide architecture	Need for scalable, resilient deployment across a city network
Operational reporting	Limited measurement of conservation and maintenance outcomes

CHAPTER 2: PROBLEM IDENTIFICATION AND ANALYSIS

2.1 Problem Identification

In the absence of a dedicated intelligent monitoring platform, most water utilities rely on a combination of SCADA systems limited to key assets, periodic manual inspection, and customer complaints to identify problems within the distribution network. A typical response to a suspected leak begins only once it becomes visible at the surface, is reported by a customer, or is inferred from a noticeable drop in reservoir levels or an unexplained increase in non-revenue water during periodic billing reconciliation. Maintenance teams are then dispatched to locate the fault manually, often using acoustic listening equipment along suspected sections of pipe, a process that can take considerable time even after a leak has been identified as present somewhere within a district.

Some utilities have adopted smart metering (Automated Meter Infrastructure, AMI) to gain finer-grained consumption data at the customer endpoint. This improves billing accuracy and can help identify unusually high consumption at individual connections, but it does not directly observe conditions within the distribution mains themselves, and therefore cannot detect or localise a leak occurring between the treatment plant and the metered connection.

Both approaches become increasingly inadequate as a network grows in size and age. Sparse SCADA instrumentation cannot observe conditions across the full extent of the distribution mains, and periodic manual inspection cannot keep pace with the rate at which faults develop across an ageing network. This motivates the analysis of existing systems in the following section.

2.2 Existing System Analysis

SCADA systems provide reliable, real-time monitoring of critical assets such as pumping stations, treatment plants, and major reservoirs, and are well established within utility operations. However, conventional SCADA deployments typically instrument only a limited number of high-value points due to the cost of wired sensor installation, leaving most of the distribution network without continuous monitoring.

Automated Meter Infrastructure (AMI) / smart metering systems provide detailed, endpoint-level consumption data and support remote meter reading, improving billing accuracy and enabling some forms of consumption anomaly detection. They do not, however, provide visibility into pressure or flow conditions within the distribution mains, and therefore cannot detect leaks occurring before water reaches the metered connection.

Manual leak detection surveys, using acoustic listening devices and correlator equipment operated by trained personnel, remain the most direct method for physically locating leaks once their approximate presence has been established. This approach is labour-intensive, requires specialised expertise, and cannot be scaled to continuous, network-wide monitoring.

Generic IoT monitoring platforms, such as open-source dashboards and time-series visualisation tools, provide flexible infrastructure for ingesting and visualising sensor data from any domain. They do not, however, include built-in leak detection logic, pipeline topology awareness, or water-domain-specific analytics, requiring substantial custom development before they can meaningfully support water distribution monitoring.

Table 2.1: Comparison of Existing Systems

System	Strengths	Limitations
SCADA (traditional)	Reliable, real-time monitoring of key assets	Sparse coverage; blind between instrumented points
AMI / Smart Metering	Accurate, detailed consumption data	No mid-network pressure/flow visibility; detects loss retrospectively
Manual leak survey	Can precisely locate confirmed leaks	Labour-intensive; not scalable; reactive
Generic IoT dashboards	Flexible; supports diverse sensor types	No domain-specific leak logic or pipeline topology awareness

Table 2.1 indicates that no single existing category of system combines dense, network-wide sensing with domain-specific leak detection analytics and predictive maintenance capability, reinforcing the gap identified in Chapter 1.

2.3 Problem Analysis

Analysis of current practices and existing systems reveals five recurring problems that directly inform the requirements of the proposed platform.

Delayed leak identification arises because most detection methods depend on a leak becoming visible, being reported by a customer, or accumulating enough loss to appear in periodic billing reconciliation, by which point substantial water has typically already been lost.

Sparse network visibility results from the cost of wired SCADA instrumentation, which restricts continuous monitoring to a small number of critical assets rather than the distribution network as a whole, leaving most pipeline segments unmonitored between inspection cycles.

Reactive rather than predictive maintenance follows from the absence of systematic failure-risk modelling; maintenance is typically performed only after a fault occurs rather than being prioritised toward pipeline segments showing early indicators of degradation.

Undetected unauthorized consumption and metering faults persist where consumption data is analysed only at the level of periodic billing rather than through continuous pattern analysis capable of flagging anomalies as they occur.

Fragmented operational visibility results from the use of separate systems, spreadsheets, and manual reports for tracking network status, maintenance activity, and water quality, making it difficult for utility managers to form a coherent, current picture of network health.

2.4 Chapter Summary

This chapter examined how water utilities currently monitor distribution networks using SCADA, smart metering, manual inspection, and generic IoT tools, and identified the inefficiencies that arise from this arrangement, including delayed leak identification, sparse network visibility, reactive maintenance, undetected anomalous consumption, and fragmented operational visibility. Analysis of existing systems confirmed that no single available approach adequately combines dense, network-wide sensing with predictive, domain-specific analytics. The resulting functional and non-functional requirements establish the justification for the system design presented in the following chapters.

CHAPTER 3: SYSTEM REQUIREMENTS SPECIFICATION (SRS)

3.1 System Requirements

This section defines the functional, non-functional, software, and hardware requirements that guide the design and implementation of the proposed platform. These requirements are derived directly from the problems and objectives identified in Chapters 1 and 2 and serve as the basis for the system design presented in the subsequent chapter.

3.1.1 Functional Requirements

Functional requirements describe the specific behaviours and capabilities that the system must provide to its users. They define what the system should do, independent of how it is technically implemented, and form the basis for the features developed during implementation. The system shall:

11. Continuously ingest pressure, flow rate, consumption, and pipeline health data from connected IoT sensors.
12. Detect leakages by identifying statistically significant deviations in pressure and flow readings.
13. Monitor water quality parameters and flag readings outside acceptable thresholds.
14. Apply predictive models to estimate infrastructure failure risk for individual pipeline segments.
15. Generate and dispatch notifications to maintenance teams when anomalies or predicted failures are detected.
16. Analyse historical and real-time consumption data to identify abnormal usage patterns.
17. Provide an operational dashboard displaying network status, active alerts, and analytics summaries.
18. Support role-based access control for administrators, control-room operators, and field maintenance staff.
19. Generate operational reports summarising water loss, maintenance response times, and conservation metrics.
20. Support remote actuation of smart valves in response to confirmed critical faults, subject to authorized operator approval.

3.1.2 Non-Functional Requirements

Non-functional requirements define the quality attributes and operational constraints of the system rather than specific behaviours. They establish standards against which the system's performance, security, and overall reliability can be evaluated.

21. Scalability: the system shall support the addition of new sensors and district-metered areas without a redesign of the core architecture.
22. Reliability and fault tolerance: the system shall continue operating and preserve collected data during partial network or component failure.
23. Security: sensor communications and user access shall be authenticated and encrypted to prevent unauthorized data access or tampering.
24. Availability: the platform shall be designed for high availability suited to continuous, city-wide operation.

25. Performance: anomaly detection and alert generation shall occur within a short, bounded interval of sensor data ingestion to support timely response.
26. Usability: the operational dashboard shall present network status and alerts in a manner accessible to non-specialist utility staff.

3.1.3 Software Requirements

Table 3.1: Software Requirements

Software	Purpose
Operating System (Linux-based server OS)	Underlying platform for development and deployment
Visual Studio Code	Primary code editor for development
Python	Backend services, data processing, and predictive model development
Node.js	Real-time API and notification services
React	Interactive front-end operational dashboard
MQTT Broker (e.g., Eclipse Mosquitto)	Lightweight messaging for sensor data transmission
Time-series Database (e.g., InfluxDB/TimescaleDB)	Storage of high-frequency sensor readings
PostgreSQL	Relational storage for users, assets, and maintenance records
scikit-learn	Development of anomaly detection and failure-prediction models
Docker	Containerized deployment of platform services
Git	Version control of the source code during development

3.1.4 Hardware Requirements

Table 3.2: Hardware Requirements

Component	Minimum Requirement
IoT Pressure Sensors	Wireless, battery- or mains-powered, suited to pipeline mounting
IoT Flow Meters	Compatible with mains and district-metered-area installation points
Water Quality Sensors	Capable of measuring turbidity, pH, and residual chlorine
Edge Gateway / Communication Module	Supports wireless transmission (e.g., LoRaWAN or cellular)
Server (development/prototype)	Quad-core processor, 16 GB RAM, 100 GB available storage

Component	Minimum Requirement
Internet Connection	Stable broadband connection for real-time data transmission

3.1.5 Chapter Summary

Defining functional, non-functional, software, and hardware requirements prior to implementation establishes a shared understanding of what the system must achieve and the constraints within which it must operate. This structured specification reduces ambiguity during development and provides a measurable basis for evaluating the completed prototype, ensuring that subsequent design decisions remain aligned with the objectives established in Chapter 1.

CHAPTER 4: SOLUTION DESIGN AND SYSTEM ANALYSIS

4.1 Introduction

This chapter presents the architectural and analytical design of the proposed platform, building on the requirements established in Chapter 3. It describes the overall system architecture, the key actors and their interactions, the data flow between major components, and the conceptual database structure, translating the requirements specification into a coherent technical design without descending into implementation-level detail such as code or configuration.

4.2 System Architecture Overview

The platform follows a layered architecture consisting of a sensor/edge layer, a communication layer, an application layer, and a data layer, a pattern well suited to distributed IoT monitoring systems. The sensor/edge layer comprises pressure sensors, flow meters, water quality sensors, and smart valves distributed across the distribution network, each connected to a lightweight edge gateway responsible for local data buffering and transmission. The communication layer uses a publish-subscribe messaging protocol (MQTT) to transmit sensor readings from edge gateways to the central platform with low overhead, suited to constrained and intermittently connected devices.

The application layer handles authentication, anomaly detection, predictive analytics, alert generation, and coordination of notifications, while the data layer separates high-frequency time-series sensor readings from durable relational records covering users, assets, maintenance history, and alert logs.

Separating time-series sensor storage from relational operational data is a deliberate architectural decision: the time-series database is optimised for high-write-throughput ingestion of frequent sensor readings, while the relational database maintains structured, query-friendly records of assets, users, and maintenance events. This separation allows each storage technology to be used according to its strengths rather than forcing a single database to serve both purposes.

4.3 Actor and Use Case Analysis

The system supports four primary actors, each corresponding to a role identified in the requirements specification: System Administrator, Control Room Operator, Field Maintenance Technician, and Data Analyst. Each actor has access to a distinct subset of system functionality, consistent with the role-based access control requirement defined in Chapter 3.

Table 4.1: Actor–Function Mapping

Actor	Primary Interactions
System Administrator	Configures sensors and district-metered areas, manages user accounts and role permissions, oversees system health
Control Room Operator	Monitors real-time dashboard, reviews alerts, approves valve actuation for confirmed faults
Field Maintenance Technician	Receives alert notifications, updates maintenance ticket status,

Actor	Primary Interactions
	records fault resolution
Data Analyst	Reviews consumption and leak analytics, generates operational reports, evaluates predictive model performance

4.4 System Data Flow

Data within the system flows through four principal processes: sensor data ingestion, anomaly detection and analysis, alert and notification generation, and maintenance ticket management.

Table 4.2: Description of Core Data Flow Processes

Process	Trigger	Outcome
Sensor data ingestion	Periodic transmission from edge gateway	Reading stored in time-series database
Anomaly detection and analysis	New reading ingested / scheduled analysis cycle	Leak, quality, or consumption anomaly flagged if thresholds or model predictions are exceeded
Alert and notification generation	Anomaly flagged	Notification dispatched to relevant maintenance team and logged
Maintenance ticket management	Technician acknowledges or resolves alert	Ticket status updated; resolution recorded for reporting and model feedback

This separation of continuous ingestion from periodic analysis allows the platform to process high-frequency sensor data while still producing discrete, actionable alerts, directly addressing the delayed leak identification problem identified in Chapter 2.

4.5 Conceptual Database Design

The data layer is organised around a combination of time-series and relational schemas, reflecting the distinct nature of high-frequency sensor readings and structured operational records. Table 4.3 summarises the principal entities identified during analysis.

Table 4.3: Principal System Entities

Entity	Description
User	Represents a registered platform user and their role
Sensor	Represents an individual IoT sensor and its installation metadata
PipelineSegment	Represents a section of the distribution network associated with one or more sensors
Reading	Represents a time-stamped sensor measurement (pressure, flow, or quality)
LeakEvent	Represents a detected or confirmed leakage event linked to a pipeline segment

Entity	Description
MaintenanceTicket	Represents a maintenance task generated from an alert and its resolution status
DistrictMeteredArea	Represents a defined zone of the network used for consumption and loss analysis
Alert	Represents a system-generated notification linked to a detected anomaly

4.6 Chapter Summary

This chapter presented the solution design derived from the requirements specification, including the layered system architecture, the actors and use cases supported by the platform, the principal data flow processes, and the conceptual database entities. These design elements collectively address the problems identified in Chapter 2, particularly delayed leak identification, sparse network visibility, and fragmented operational data, and provide the architectural foundation for implementation.

CHAPTER 5: RESULTS AND RECOMMENDATIONS

5.1 Introduction

This chapter presents the outcomes of the capstone project relative to the objectives defined in Chapter 1, evaluates the extent to which the identified problems have been addressed, and offers recommendations for future development beyond the scope of this report.

5.2 Summary of Results

The proposed design successfully addresses the core problems identified in Chapter 2. Table 5.1 summarises the project objectives against their corresponding outcomes.

Table 5.1: Objectives and Corresponding Outcomes

Objective	Outcome
Continuous sensor data ingestion	Achieved through an MQTT-based communication layer feeding a time-series database
Leak detection via pressure/flow anomaly analysis	Achieved through statistical threshold and predictive-model-based anomaly detection
Water quality monitoring	Achieved through dedicated quality sensor ingestion and threshold-based flagging
Predictive infrastructure failure modelling	Achieved at design level using historical and sensor-derived features
Alerting and notification mechanism	Achieved through automated alert generation linked to maintenance ticketing
Consumption pattern analysis	Achieved through comparison of real-time consumption against historical baselines
Operational dashboard	Achieved through a React-based interface summarising network status and alerts
Role-based access control	Achieved through defined actor roles with differentiated permissions

5.3 Evaluation Against Requirements

The functional and non-functional requirements defined in Chapter 3 were used as the primary evaluation criteria. Core functional requirements relating to data ingestion, leak detection, alerting, and role-based access were satisfied by the design presented in Chapter 4. Non-functional requirements, particularly scalability and fault tolerance, are dependent on the underlying messaging and storage infrastructure's behaviour under sustained load and were considered adequately addressed at the design level, though full validation requires the performance testing typically conducted during a later implementation and testing phase.

5.4 Limitations

As an academic capstone project, the system has scope-related limitations rather than implementation defects at this design stage. These include the exclusion of full city-wide deployment, reliance on simulated or limited-scale sensor data during prototyping, and a predictive maintenance model whose accuracy has not yet been validated against extensive historical failure records. The smart-valve actuation feature is limited to a supervised, approval-gated action rather than a fully autonomous control loop, consistent with the safety-conscious scope of an academic prototype.

5.5 Chapter Summary

The project met its stated objectives at the design and prototype level, demonstrating that continuous IoT-based monitoring, leak detection, and predictive maintenance can be integrated using established messaging and analytics frameworks. Remaining limitations are consistent with the scope defined for an academic capstone project, and the recommendations outlined above indicate a clear path for extending the platform toward a more complete system.

CHAPTER 6: REFLECTION ON LEARNING AND PERSONAL DEVELOPMENT

6.1 Introduction

This chapter reflects on the technical and professional learning gained throughout the development of this capstone project. It considers the skills acquired, the challenges encountered, and the areas identified for continued growth as a software engineer.

6.2 Technical Learning

Working on this project required engagement with concepts that extended beyond prior coursework, particularly around IoT messaging protocols and time-series data management. Researching MQTT-based communication and time-series database design provided a practical understanding of how high-frequency sensor data can be ingested and stored efficiently, a topic that had previously been understood only in general terms. Designing the anomaly detection approach in Chapter 4 also reinforced the importance of separating transient, high-volume sensor data from durable operational records, a distinction that was not immediately obvious at the outset of the project.

Working through the requirements analysis in Chapters 2 and 3 was equally instructive. Translating a broad and somewhat ambiguous problem, that of inefficient water loss detection, into specific, testable functional and non-functional requirements demonstrated the value of structured requirements engineering over informal feature planning.

6.3 Challenges Encountered

The most significant challenge was scoping the project appropriately for an academic timeframe. The initial feature set, encompassing real-time monitoring, leak detection, predictive maintenance, water quality analysis, and city-wide fault tolerance, was broad enough to resemble a full utility-management product, and considerable effort went into distinguishing which aspects were essential to demonstrating the core technical contribution and which were reasonable to simplify or exclude.

A further challenge was reasoning about a domain, water distribution engineering, that was not part of prior technical training. Understanding why pressure and flow anomalies indicate leakage, and how non-revenue water is measured and reported by utilities, required additional research and highlighted the importance of domain analysis before committing to a technical design.

6.4 Development of Professional Skills

Beyond technical knowledge, the project contributed to skills relevant to professional software engineering practice. Structuring the report itself, moving systematically from problem identification through requirements to design, mirrored the documentation discipline expected in industry settings, where design decisions must be justified and traceable to stated requirements rather than presented as isolated technical choices. Evaluating existing systems in Chapter 2 also developed a more critical approach to assessing available tools, rather than assuming that a new system is warranted without first establishing what existing solutions fail to provide.

6.5 Areas for Further Development

This project identified areas warranting further development, both technically and academically. Practical experience implementing and testing the anomaly detection and predictive maintenance models against real or realistically simulated sensor datasets, rather than analysing them at a design level, would deepen understanding of their performance and reliability characteristics. Additionally, exposure to field trials with actual utility sensor deployments would strengthen the ability to validate design assumptions against real network behaviour rather than relying solely on requirements analysis.

CHAPTER 7: CONCLUSION

7.1 Summary of the Study

This report examined the problem of inefficient water loss detection and infrastructure monitoring in urban distribution networks, a domain in which utilities continue to rely on sparse SCADA instrumentation, periodic manual inspection, and customer complaints to identify problems that, by the time they are detected, have often already resulted in substantial water loss. Chapter 1 established the background of water distribution monitoring and the industry context within which IoT-based sensing and analytics have developed, noting that while general-purpose IoT platforms have matured considerably, their adoption for dense, network-wide water monitoring remains limited. Chapter 2 extended this discussion through a detailed problem analysis, showing that existing systems, whether traditional SCADA, smart metering, or generic IoT dashboards, fail to combine dense sensor coverage with domain-specific leak detection and predictive maintenance analytics. This gap was shown to produce specific, recurring problems: delayed leak identification, sparse network visibility, reactive rather than predictive maintenance, undetected anomalous consumption, and fragmented operational visibility.

In response to these findings, this report proposed and designed an IoT-enabled water distribution monitoring platform built on established messaging and time-series storage technologies, rather than a custom-built telemetry stack. This design decision allowed the project to focus its engineering effort on the domain-specific aspects of water distribution monitoring, including leak detection logic, predictive failure modelling, and role-differentiated operational workflows, rather than on reimplementing sensor communication and data-ingestion mechanisms that are already well established in existing open-source frameworks.

7.2 Achievement of Objectives

The objectives defined in Section 1.7 were addressed systematically across the report. The requirements analysis presented in Chapters 2 and 3 translated the broad aim of the project into a concrete set of functional and non-functional requirements, ranging from continuous sensor data ingestion and leak detection to role-based access control and operational reporting. These requirements were then mapped directly onto the architectural and analytical design presented in Chapter 4, which described a layered system architecture separating sensor, communication, application, and data layers, an actor-based model reflecting the distinct responsibilities of administrators, control-room operators, field technicians, and data analysts, and a conceptual database structure capable of supporting both high-frequency sensor storage and durable operational records.

Chapter 5 evaluated the resulting design against the original requirements and confirmed that the core objectives, particularly continuous monitoring, leak detection, and alerting, were achieved at the design and prototype level. Objectives related to supplementary functionality, such as predictive infrastructure failure modelling and operational analytics, were also incorporated into the design, though their depth was appropriately scoped to remain consistent with the constraints of an academic capstone project.

7.3 Contribution of the Project

The primary contribution of this project lies in demonstrating a practical architectural approach to a gap identified between sparse, asset-focused SCADA monitoring and dense, network-wide sensing supported by modern IoT technologies. Rather than treating continuous monitoring and domain-specific leak analytics as separate concerns, as appeared to be the case across the existing systems evaluated in Chapter 2, this project shows that the two can be integrated by building water-domain-specific analytics on top of established IoT messaging and time-series storage frameworks. This approach allows a relatively focused development effort to achieve monitoring and predictive-maintenance capabilities that would otherwise require substantial investment in custom telemetry infrastructure.

Beyond its technical contribution, the project also illustrates the value of structured requirements engineering and architectural justification in addressing a problem that, at first appearance, might be treated as a straightforward instrumentation gap. By tracing each design decision back to a specific requirement and, in turn, to a specific problem identified through analysis of existing utility practices, the report demonstrates a disciplined approach to software engineering practice that extends beyond the immediate technical implementation.

7.4 Limitations and Future Direction

As discussed in Chapter 6, the system remains a prototype scoped deliberately for academic purposes, and several aspects of the design, particularly predictive model accuracy under real-world failure data, performance under city-wide sensor density, and field usability with actual utility personnel, were addressed at the design level rather than validated through extensive empirical testing. These limitations are consistent with the scope defined for the project and do not diminish the validity of the design, but they do delineate the boundary between what this report has established and what remains for future work. Recommendations for extending the platform include field trials with real sensor deployments, refinement of predictive models using verified historical failure data, and formal usability evaluation with utility control-room staff.

7.5 Final Remarks

Taken as a whole, this report has moved systematically from the identification of a genuine gap in existing water distribution monitoring practices, through structured requirements analysis, to a justified architectural design intended to close that gap. While the system described remains a prototype rather than a production-ready application, the analysis, design, and evaluation presented across the preceding chapters establish a coherent and technically grounded foundation for an intelligent water distribution monitoring platform. With further development, empirical testing, and refinement along the lines suggested in Chapter 6, the underlying design demonstrated in this report could reasonably be extended toward a more complete system suited to real-world municipal water management.

REFERENCES

- H. Kadar; et al., "SMART2L: Smart Water Level and Leakage Detection", 2018
- S. Sarangi, "Smart Water Leakage and Theft Detection using IoT", 2020
- K. B. Adedeji; A. M. Abu-Mahfouz, "Towards Achieving a Reliable Leakage Detection and Localization Algorithm for Application in Water Piping Networks: An Overview", 2023
- S. Thenmozhi; et al., "IoT Based Smart Water Leak Detection System for a Sustainable Future", 2021
- M. M. Ali; et al., "IoT-Driven Water Leak Detection: A Comparative Study of Machine Learning and Deep Learning Models for Time-Series Sensor Data with Explainable AI", 2024
- Y. Asadi, "Employing Machine Learning in Water Infrastructure Management: Predicting Pipeline Failures for Improved Maintenance and Sustainable Operations", Industrial Artificial Intelligence, 2(1), 2024
- R. Chen; Q. Wang; A. Javanmardi, "A Review of the Application of Machine Learning for Pipeline Integrity Predictive Analysis in Water Distribution Networks", Archives of Computational Methods in Engineering, 32(6), 2025
- K. Kakoudakis; R. Farmani; D. Butler, "Pipeline Failure Prediction in Water Distribution Networks Using Weather Conditions as Explanatory Factors", Journal of Hydroinformatics, 20, 2018
- N. A. Barton; T. S. Farewell; S. H. Hallett; et al., "Improving Pipe Failure Predictions: Factors Affecting Pipe Failure in Drinking Water Networks", Water Research, 164, 2019